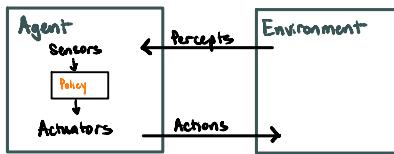


## 1 Introduction

- An agent perceives its environment through sensors and acts upon it through actuators
- A rational agent chooses action to max expected value



## 2 Search

Search Problems have:

- State Space  $S$
- Initial State  $s_0$
- Actions in state  $A(s)$

- Transition model  $\text{Result}(s, a)$
- Goal test  $G(s)$
- Action cost  $c(s, a, s')$

Solutions: action sequence that reaches goal state

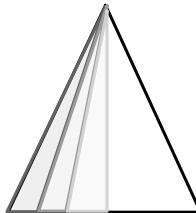
Optimal Solution: least cost among solutions

### Depth First Search (DFS)

Strategy: expand a deepest node first

Use: LIFO stack

Optimal?: No, finds "leftmost" solution

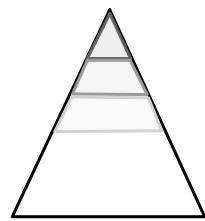


### Breadth First Search (BFS)

Strategy: expand shallowest node first

Use: FIFO Queue

Optimal: If costs are equal

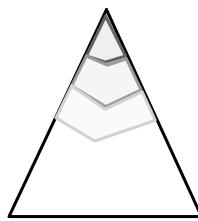


### Uniform Cost Search (UCS)

Strategy: expand lowest cost from root  $g(n)$

Use: Priority Queue sorted by  $g(n)$

Optimal?: Yes

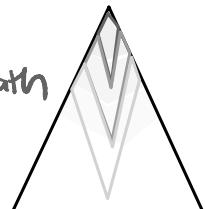


### A\* Search

Strategy: expand node most likely to be on optimal path

Use: Priority Queue sorted by  $g(n) + h(n)$

Optimal: Yes, if admissible and consistent



Admissible heuristic:  $0 \leq h(n) \leq h^*(n)$   
heuristic  $\leq$  actual cost ↑ true cost to goal

Consistent heuristic:  $h(A) - h(C) \leq c(A, C)$

heuristic "arc" cost  $\leq$  actual cost for each arc

## Greedy Search

Strategy: selects frontier node w/ lowest heuristic value  
for expansion, becomes closest to goal

Use: priority queue with heuristics, estimated forward cost

Optimal: No, not guaranteed to find goal state, unpredictable

## Local Search

### Hill Climbing

Idea: Start wherever, repeat: move to the best neighboring state  
If no neighbors better than current, quit

function HILL-CLIMBING(problem) returns state  
current  $\leftarrow$  make-node(problem.initial state)

loop do

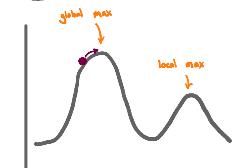
neighbor  $\leftarrow$  highest valued successor of current

if neighbor.value  $\leq$  current.value then

return current.state

current  $\leftarrow$  neighbor

Random-restart hill climbing trivially complete: restarts from random initial state



### Simulated Annealing

Idea: random walk + hill climbing, choose randomly, choose worse w/ some probability according to temperature, temp starts high w/ more "bad" moves allowed and decreases

```

function SIMULATED_ANNEALING(problem, schedule) returns state
    current ← problem.initial-state
    for t=1 to ∞ do
        T ← schedule(t)
        if T = 0 then return current
        next ← a randomly selected successor of current
        ΔE ← next.value - current.value
        if ΔE > 0 then current ← next
        else current ← next w/ prob  $e^{\frac{-\Delta E}{T}}$ 

```

If  $T$  decreased slowly enough, will converge to optimal state

## Local Beam Search

multiple interaction searches

Idea: track  $k$  states at each iteration,  $k$  threads share information  
 good threads attract other threads, chooses  $k$  best successors

## Genetic Algorithms

Break and recombine states

Idea: beam search  $k$  states in population, each state evaluated  
 w/ evaluation function (fitness func), offspring produced by  
 crossing parent strings at crossover point

## (3) Games

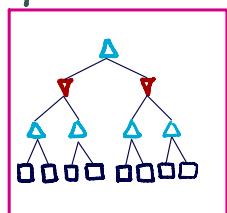
Zero sum games: our gain is directly equivalent to opponent's loss

### Minimax

algorithm to counter optimal opponent makes

Idea: zero sum game algorithm assuming opponent plays  
 optimally

terminal utilities/state value: optimal score attainable



function minimax-decision(s) returns action

return action  $a$  in Actions(s) w/ highest minimax-value (Result(s,a))

function minimax-value(s) returns a value

if Terminal-Test(s) then return Utility(s)

if Player(s)=MAX then return  $\max_{a \in \text{Actions}(s)} \text{minimax-value}(\text{Result}(s,a))$

if Player(s)=MIN then return  $\min_{a \in \text{Actions}(s)} \text{minimax-value}(\text{Result}(s,a))$

## Alpha-Beta Pruning

optimization to minimax:  $O(b^n) \Rightarrow O(b^{n/2})$

Idea: Stop looking as soon as you know n's value can at best equal optimal value of n's parent

$\alpha$ : MAX's best option on path to root

$\beta$ : MIN's best option on path to root

def max-value(state,  $\alpha$ ,  $\beta$ )

initialize  $V = -\infty$

for each successor of state:

$V = \max(V, \text{value}(\text{successor}, \alpha, \beta))$

if  $V \geq \beta$  return  $V$

$\alpha = \max(\alpha, V)$

return  $V$

def min-value(state,  $\alpha$ ,  $\beta$ )

initialize  $V = +\infty$

for each successor of state:

$V = \min(V, \text{value}(\text{successor}, \alpha, \beta))$

if  $V \leq \alpha$  return  $V$

$\beta = \min(\beta, V)$

return  $V$

Evaluation functions are used in depth-limited minimax to output an estimate of the true minimax value of the node

$$\text{Eval}(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

Weight associated w/ feature

features of state

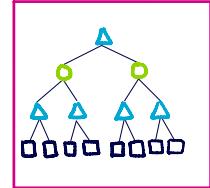
## Expectimax

algorithm with randomness using Expected value

Idea: Introduce chance nodes calculated by taking expected utility of children

function decision(s) returns action

return action  $a$  in Actions(s) w/ highest value (Result(s,a))



function value(s) returns a value

if Terminal-Test(s) then return Utility(s)

if Player(s)= MAX then return  $\max_{a \in \text{Actions}(s)} \text{value}(\text{Result}(s,a))$

if Player(s)= MIN then return  $\min_{a \in \text{Actions}(s)} \text{value}(\text{Result}(s,a))$

if Player(s)= CHANCE then return  $\sum_{a \in \text{Actions}(s)} \text{Pr}(a) * \text{value}(\text{Result}(s,a))$

## Monte Carlo Tree Search (MCTS)

Idea:

1) Evaluation by rollouts: From state  $s$  play many times using policy and count wins and losses

2) Selective search: explore parts of the tree, without constraints on horizon, that will improve decision at root

## UCB Algorithm

$$\text{UCB1}(n) = \frac{U(n)}{N(n)} + C$$

total wins for Player(Parent(n))

$\sqrt{\frac{\log N(\text{PARENT}(n))}{N(n)}}$

total rollouts from  $n$

User specific param to balance exploration vs. exploitation

## 4 Logic

Maintain a knowledge base w/ logical sentences that can make logical inferences

## Logic Syntax

Symbol	Meaning	Description
$\neg$	not	$\neg P$ true iff $P$ is false
$\wedge$	and (conjunction)	$A \wedge B$ true iff both $A$ true and $B$ true
$\vee$	or (disjunction)	$A \vee B$ true iff either $A$ true or $B$ true
$\Rightarrow$	implication	$A \Rightarrow B$ true unless $A$ true and $B$ false
$\Leftrightarrow \equiv$	biconditional	$A \Leftrightarrow B$ ( $A \equiv B$ ) true iff either both $A, B$ true or false

**Conjunctive normal form (CNF):** conjunction of clauses which are disjunction of literals

$$(P_1 \vee \dots \vee P_i) \wedge \dots \wedge (P_j \vee \dots \vee P_n)$$

Converting to CNF:

- 1) Eliminate  $\Leftrightarrow$
- 2) Eliminate  $\Rightarrow$
- 3) Not's ( $\neg$ ) only on literals
- 4) Reduce and distribute

$$\begin{aligned} (\alpha \Leftrightarrow \beta) &\equiv ((\alpha \Rightarrow \beta) \wedge (\beta \Rightarrow \alpha)) \\ \alpha \Rightarrow \beta &\equiv \neg \alpha \vee \beta \\ \neg(\alpha \wedge \beta) &\equiv (\neg \alpha \vee \neg \beta), \neg(\alpha \vee \beta) \equiv (\neg \alpha \wedge \neg \beta) \\ (\alpha \wedge (\beta \vee \gamma)) &\equiv ((\alpha \wedge \beta) \vee (\alpha \wedge \gamma)), \\ (\alpha \vee (\beta \wedge \gamma)) &\equiv ((\alpha \vee \beta) \wedge (\alpha \vee \gamma)) \end{aligned}$$

**model:** assignment of true/false to all proposition symbols

**valid:** true in all models

**satisfiable:** at least one model where true

**unsatisfiable:** not true in any models

## First Order Logic (FOL)

uses objects as base components

use function symbols to name objects in terms  
quantifiers:

universal quantifier  $\forall$ : "for all"  
existential quantifier  $\exists$ : "there exists"

**Operator Precedence**  
 $\neg, =, \wedge, \vee, \Rightarrow, \Leftrightarrow$

## Propositional Logical Inference

**Entails ( $\models$ ):** in all models where  $A$  is true,  $B$  is as well,  $M(A) \subseteq M(B)$

Show by:

- 1)  $A \models B$  iff  $A \Rightarrow B$  valid
- 2)  $A \models B$  iff  $A \wedge \neg B$  unsatisfiable

direct proof

proof by contradiction

## Model Checking

### DPLL Algorithm

depth-first, backtracking search w/ tricks SAT solver

Idea: Solve satisfiability problem given in CNF, continue assigning truth values until found or cannot be found, backtrack w/ conditions:

1) Early Termination: clause true if any symbol true, sentence false if any single clause is false

2) Pure Symbol Heuristic: symbol only shows up in pos/neg form assigned immediately  $(A \vee B)(\neg B \vee C) \wedge (\neg C \vee A)$   $A = \text{true}$

3) Unit Clause Heuristic: clause w/ one literal or literal w/ many falses

## Theorem Proving

### Resolution Algorithm

Idea: iteratively applies it to knowledge base either q inferred or nothing left to infer

### Forward Chaining Algorithm

Idea: data driven reasoning, iterating through every implication statement where LHS known, adding RHS to facts using generalized Modus Ponens

Solve inference in FOL by generalized Modus Ponens or propositionalization, translating problem into propositional logic and using SAT solver

5

## Probability

Marginal distribution:

$$P(A, B) = \sum_C P(A, B, C)$$

Bayes' rule:

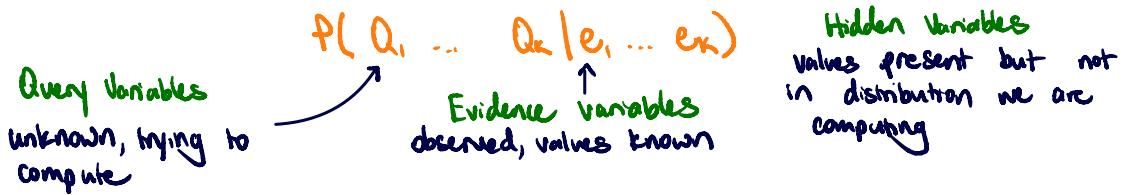
$$P(A|B, C) = \frac{P(A, B, C)}{P(B, C)} = \frac{P(A, B | C)}{P(B | C)}$$

$$\text{Chain Rule: } P(A, B, C) = P(A, B|C)P(C) = P(A|B, C)P(B|C)P(C)$$

$$\text{A conditionally independent given } C: \quad P(A, B|C) = P(A|C)P(B|C)$$

$$\text{A independent of } B \text{ given } C: \quad P(A|B, C) = P(A|C)$$

$$\text{A, B independent:} \quad P(A, B) = P(A)P(B)$$



## Bayes Nets

Idea: Joint distribution distributed across smaller probability tables with Directed Acyclic Graph (DAG)

- 1) DAG w/ node per variable  $X$
- 2) conditional distribution for each node  $P(X|A_1, \dots, A_n)$  where  $A_i$  is  $i^{\text{th}}$  parent of  $X$ , stored as conditional probability table (CPT)

Each node is conditionally independent of all ancestor nodes in graph, given all of parents

## Bayes Net Inference (calculating joint probability)

Eliminate variables by:

1. Join (multiply together) all factors involving  $X$
2. Sum out  $X$

## Prior Sampling

randomly generate samples

## Rejection Sampling

early reject any sample inconsistent w/ evidence

## Likelihood Weighting

Ensure we never generate a bad sample

- 1) manually set variables to equal evidence
- 2) weight each sample by probability of the evidence variables given the sampled variables

## Gibbs Sampling

Set all variables to random value, repeatedly pick one variable at a time, clear value, resample

## 6 Markov Models



Chain like infinite-length Bayes Net

Need to know initial state and transition model

## Mini-Forward Algorithm

$$\Pr(W_{i+1}) = \sum_{w_i} \Pr(W_{i+1} | w_i) \Pr(w_i)$$

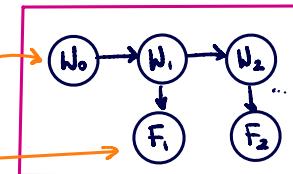
## Stationary Distribution

$$\Pr(W_{t+1}) = \Pr(W_t) = \sum_{w_t} \Pr(W_{t+1} | w_t) \Pr(w_t)$$

# Hidden Markov Models

allows to observe some evidence at timestep  
and affects belief distribution

State variable  
encodes belief  
evidence var



## Forward Algorithm

$$B'(W_{t+1}) = \sum_{w_t} \Pr(W_{t+1} | w_t) B(w_t)$$

Time elapse update

$$B(W_{t+1}) \propto \Pr(f_{t+1} | W_{t+1}) B'(W_{t+1})$$

Observation update

$$B(W_{t+1}) \propto \Pr(f_{t+1} | W_{t+1}) \sum_{w_t} \Pr(W_{t+1} | w_t) B(w_t)$$

## Viterbi Algorithm

Dynamic Programming algorithm to solve

$$\operatorname{argmax}_{x_{1:N}} P(x_{1:N} | e_{1:N}) = \operatorname{argmax}_{x_{1:N}} P(x_{1:N}, e_{1:N})$$

$$m_r[x_r] = P(e_r | x_r) m_{r-1} P(x_r | x_{r-1}) m_{r-1}[x_{r-1}]$$